

Proposal slots

matched ask count

filter worse

−0.83 pp terminal
@ 32,500 slots

Real simulator evaluations

matched run_world calls

filter better

+3.65 pp coverage
@ 20,000 evals

Wall / LLM calls

clock time, not invoices

filter slower

~20 vs ~7 min
same skip gate